


Matrix multiplication and dense neural networks

Prof. Luca Formaggia, Matteo Caldana

September 22, 2023



1 Introduction

The simple matrix-matrix multiplication is the operation that stands at the basis of almost all neural networks, however an efficient implementation is by no means trivial.

2 Objectives

Write a code that performs the matrix-matrix multiplication [1]. Your code should support both `float` and `double`. The size of the matrix should be able to be specified at run-time. Run strong and weak scalability test and report the results. Compare the results and elapsed time of your code with the ones of the OpenBLAS library (available in the `mk` modules) [2]. Employ the following techniques to make your code go faster, quantify the improvements you get with each of them:

- Exploit the SIMD instruction [3].
- Exploit loop unrolling [4].
- Make sure that your code is cache friendly [5,6,7].
- Use OpenMP to parallelize the code.

Use the matrix-matrix multiplication to build and train a dense feed-forward neural network. You can use a simple dataset of your choice and hard-code the computations for the gradient and use a simple gradient descent as optimizer.

3 Ideas for an exam project

- Parallelize the code with CUDA.
- Make your code more general: automatize the computations of the gradient, handle different activation functions, implement more sophisticated optimizers (such as momentum or ADAM).

References

- [1] [https://inst.eecs.berkeley.edu/~cs61c/fa17/lec/18/L18%20SIMD%20\(1up\).pdf](https://inst.eecs.berkeley.edu/~cs61c/fa17/lec/18/L18%20SIMD%20(1up).pdf)
- [2] <https://github.com/OpenMathLib/OpenBLAS>
- [3] https://en.wikipedia.org/wiki/Single_instruction,_multiple_data
- [4] https://en.wikipedia.org/wiki/Loop_unrolling
- [5] https://en.wikipedia.org/wiki/Cache_performance_measurement_and_metric
- [6] https://en.wikipedia.org/wiki/Cache_hierarchy
- [7] <https://siboehm.com/articles/22/Fast-MMM-on-CPU>